


# Warmup #2

Bill Cheng

*<http://merlot.usc.edu/cs402-f18>*



# Multi-threading Exercise



Make sure you are familiar with the *pthread*s library

– Ch 2 of textbook - threads, signals

- additional resource is a book by Nichols, Buttlar, and Farrell  
“*Pthreads Programming*”, O’Rielly & Associates, 1996

– you must learn how to use pthreads *mutex* and *condition variables* correctly


- `pthread_mutex_lock()` / `pthread_mutex_unlock()`
- `pthread_cond_wait()` / `pthread_cond_signal()` / `pthread_cond_broadcast()`

– you must learn how to handle UNIX *signals* (<Cntrl+C>)


- `pthread_sigmask()` / `sigwait()`
- `pthread_kill()` / `pthread_cancel()`

– you need to learn how to perform *cancellation* in pthreads

- `pthread_setcancelstate()`
- `pthread_setcanceltype()`
- `pthread_testcancel()`




# Token Bucket Filter





**Ex:**

- ▬ ticket scalper?!
- ▬ traffic controller/shaper




# Arrivals & Departures

- ▢  $a_i$  : arrival time
- ▢  $d_i$  : departure time
- ▢  $s_i$  : service time
- ▢  $r_i$  : response (system) time
- ▢  $q_i$  : queueing/waiting time





▢  $r_1 = d_1 - a_1$




# Arrivals & Departures

- ▢  $a_i$  : arrival time
- ▢  $d_i$  : departure time
- ▢  $s_i$  : service time
- ▢  $r_i$  : response (system) time
- ▢  $q_i$  : queueing/waiting time





▢  $r_2 = d_2 - a_2$




# Arrivals & Departures

- ▢  $a_i$  : arrival time
- ▢  $d_i$  : departure time
- ▢  $s_i$  : service time
- ▢  $r_i$  : response (system) time
- ▢  $q_i$  : queueing/waiting time




▢  $r_3 = d_3 - a_3$





# Event Driven Simulation


- ➡ An **event queue** is a sorted list of events according to timestamps; smallest timestamp at the head of queue
- ➡ **Object oriented**: every object has a "next event" (what it will do next if there is no interference), this event is inserted into the event queue
- ➡ Execution: remove an event from the head of queue, "execute" the event (notify the corresponding object so it can insert the next event)
- ➡ Insert into the event queue according to timestamp of a new event; insertion may cause additional events to be deleted or inserted
- ➡ Potentially repeatable runs (if the same seed is used to initialize random number generator)



# Event Driven Simulation (Cont...)




$$r_3 = d_3 - a_3$$





# Time Driven Simulation

- ➡ Every active object is a thread
- ➡ To execute a job for  $x$  msec, the thread sleeps for  $x$  msec
  - nunki.usc.edu does not run a realtime OS
  - it may not get woken up more than  $x$  msec later, and sometimes, *a lot more* than  $x$  msec later
    - you need to decide if the extra delay is reasonable or it is due to a bug in your code
- ➡ Let your machine decide which thread to run next (irreproducible results)
- ➡ Compete for resources (such as Q1), must use mutex




# Time Driven Simulation (Cont...)


➡ You will need to implement 4 threads (or 1 main thread and 4 child threads)

— the *packet arrival thread* sits in a loop

- sleeps for an interval, trying to match a given *inter-arrival time* (from trace or deterministic)
- wakes up, creates a packet object, locks mutex
- enqueues the packet to Q1
- moves the first packet in Q1 to Q2 if there are enough tokens
- if Q2 was empty before, need to *signal* or *broadcast* a *queue-not-empty condition*
- unlocks mutex
- goes back to sleep for the "right" amount



➡ See skeleton code in warmup2 FAQ

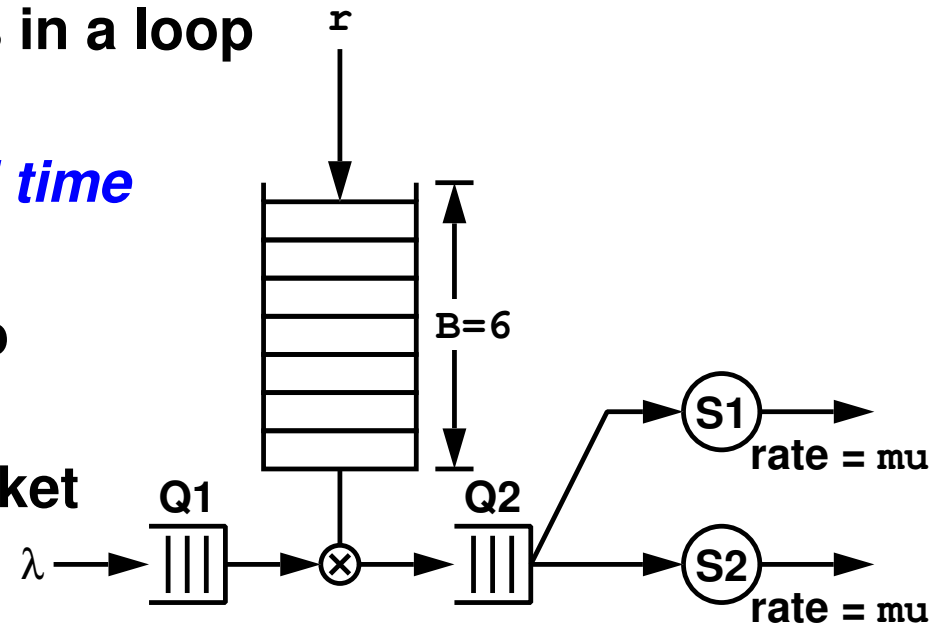


# Time Driven Simulation (Cont...)


➡ You will need to implement 4 threads (or 1 main thread and 4 child threads)

— the *token depositing thread* sits in a loop

- sleeps for an interval, trying to match a given *inter-arrival time* for tokens
- wakes up, locks mutex, try to increment token count
- check if it can move first packet from Q1 to Q2
- if packet is added to Q2 and Q2 was empty before, *signal* or *broadcast* a *queue-not-empty condition*
- unlocks mutex
- goes back to sleep for the "right" amount



➡ See skeleton code in warmup2 FAQ



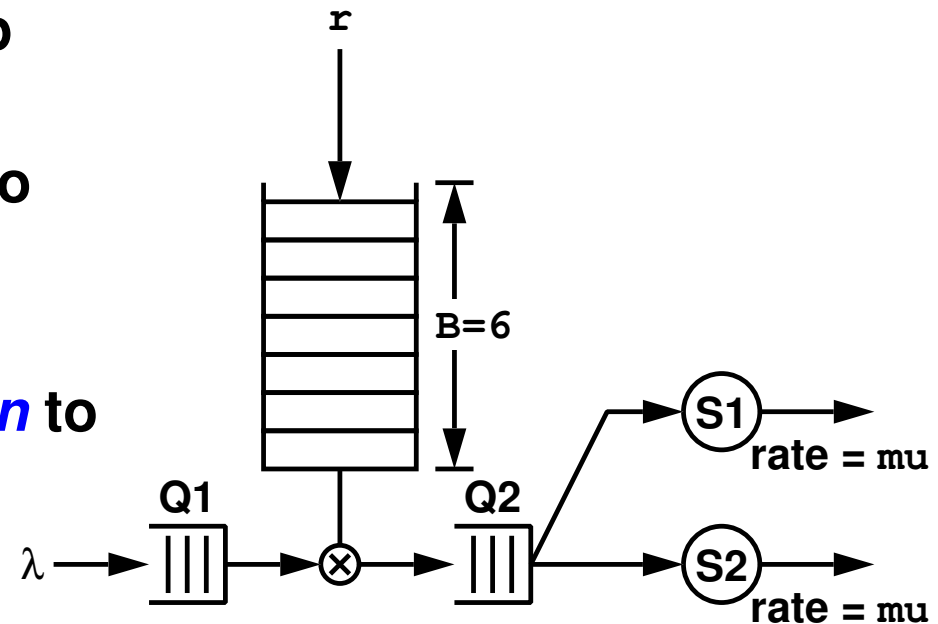
# Time Driven Simulation (Cont...)

➡ You will need to implement 4 threads (or 1 main thread and 4 child threads)


— each *server thread* sits in a loop

- wait for work {
- lock mutex
  - stay in a loop to wait for Q2 to become non-empty
    - ◇ if empty, *wait* for the *queue-not-empty condition* to be *signaled/broadcasted*
  - when Q2 is not empty, dequeue a packet
  - unlock mutex

- work {
- sleeps for an interval matching the *service time* of the packet
  - eject the packet from the system
  - loop




➡ See skeleton code in warmup2 FAQ



# Time Driven Simulation (Cont...)

- ➡ **Dropped packets**
  - if the token requirement for an arriving packet is too large, drop the packet
- ➡ **Dropped tokens**
  - if an arriving token finds a full bucket, it is dropped
- ➡ **Other requirements**
  - please read the spec!



# Program Output

- ➡ Program output must look like what's in the spec
- you must **NOT** wait for emulation to end to print all these

## Emulation Parameters:

```


number to arrive = 20
lambda = 2          (if -t is not specified)
mu = 0.35           (if -t is not specified)
r = 4
B = 10
P = 3               (if -t is not specified)
tsfile = FILENAME   (if -t is specified)

```

```

00000000.000ms: emulation begins
00000251.726ms: token t1 arrives, token bucket now has 1 token
00000502.031ms: token t2 arrives, token bucket now has 2 tokens
00000503.112ms: p1 arrives, needs 3 tokens, inter-arrival time = 503.112ms
00000503.376ms: p1 enters Q1
00000751.148ms: token t3 arrives, token bucket now has 3 tokens
00000751.186ms: p1 leaves Q1, time in Q1 = 247.810ms, token bucket now has 0 token
00000752.716ms: p1 enters Q2
00000752.932ms: p1 leaves Q2, time in Q2 = 0.216ms
00000752.982ms: p1 begins service at S1, requesting 2850ms of service
00001004.271ms: p2 arrives, needs 3 tokens, inter-arrival time = 501.159ms
00001004.526ms: p2 enters Q1
00001007.615ms: token t4 arrives, token bucket now has 1 token
00001251.259ms: token t5 arrives, token bucket now has 2 tokens
00001505.986ms: p3 arrives, needs 3 tokens, inter-arrival time = 501.715ms
00001506.713ms: p3 enters Q1
00001507.552ms: token t6 arrives, token bucket now has 3 tokens
00001508.281ms: p2 leaves Q1, time in Q1 = 503.755ms, token bucket now has 0 token
00001508.761ms: p2 enters Q2
00001508.874ms: p2 leaves Q2, time in Q2 = 0.113ms
00001508.895ms: p2 begins service at S2, requesting 1900ms of service
...

```



# Program Output

- ➡ Program output must look like what's in the spec
- you must **NOT** wait for emulation to end to print all these

```
...
00003427.557ms: p2 departs from S2, service time = 1918.662ms, time in system = 2423.286ms
00003612.843ms: p1 departs from S1, service time = 2859.861ms, time in system = 3109.731ms
...
?????????.???ms: p20 departs from S?, service time = ????.???ms, time in system =
?????.???ms
?????????.???ms: emulation ends
```


## Statistics:

```
average packet inter-arrival time = <real-value>
average packet service time = <real-value>
```


```
average number of packets in Q1 = <real-value>
average number of packets in Q2 = <real-value>
average number of packets at S1 = <real-value>
average number of packets at S2 = <real-value>
```

```
average time a packet spent in system = <real-value>
standard deviation for time spent in system = <real-value>
```

```
token drop probability = <real-value>
packet drop probability = <real-value>
```



# Calculating Statistics



**NOTE:** this is just one scenario

- other scenarios are possible
- events on right may **not** be correct

arrival thread timeout (read clock?)

should you print arrival msg here?

call to lock mutex to enter Q1

enter Q1 (read clock)

leave Q1 (read clock)

remove token(s)

enter Q2 (read clock)

leave Q2 (read clock)

unlock mutex

should you print leave queue and begin service msgs here?

begin service at S1

leave S1

should you print leave msg here?

overhead?

time in Q1

time in Q2

time in S1


charge to no one

time

- time between **begin service** and **leave server** is the amount of time in `select()` or `usleep()`

➡ Some packets need to be excluded from certain statistics

- e.g., if a packet is dropped, it should not participate in the **time-in-system** statistics







# Mean and Standard Deviation

- ➡ **Average time**  
 — for  $n$  samples, add up all the time and divide by  $n$


- ➡ **Average number of packets at a server**  
 — same a fraction of time the server is busy



- ➡ **Average number of packets at Q1**



- ➡ **Standard deviation is the squareroot of variance**  
 —  $Var[X] = E[X^2] - (E[X])^2$   
 ○ must use *population variance* equation



# SIGINT



<Cntrl+C>

- you need to *termiante* your emulation *gracefully* in case the user presses <Cntrl+C>
  - the OS will *delivery a SIGINT Unix signal to your process*
  - you need to *catch the signal*
  - you need to let all your threads know that it's "time to quit"
  - you need to *wait* for all your threads to terminate
  - then you can print statistics and quit
- ➡ We will cover this in lecture next week
  - will discuss this at the beginning of the next discussion section

